Same 64 operations, fewer compiled regions

Actual generated ATen graph fragments · depth 8, seed 0 · PyTorch 2.14.1+cu130 on B200

False: 16 regions

AOT 0 4 ATen ops	**AOT 1** 4 ATen ops	**AOT 2** 4 ATen ops	**AOT 3** 4 ATen ops
AOT 4 4 ATen ops	**AOT 5** 4 ATen ops	**AOT 6** 4 ATen ops	**AOT 7** 4 ATen ops
AOT 8 4 ATen ops	**AOT 9** 4 ATen ops	**AOT 10** 4 ATen ops	**AOT 11** 4 ATen ops
AOT 12 4 ATen ops	**AOT 13** 4 ATen ops	**AOT 14** 4 ATen ops	**AOT 15** 4 ATen ops

16 replay launches + 16 copy kernels per call

True: 2 regions

AOT 0: 32 ATen ops

sin → multiply → add residual → cos

8 repeated blocks fused into one kernel

AOT 1: 32 ATen ops

sin → multiply → add residual → cos

8 repeated blocks fused into one kernel

2 replay launches + 2 copy kernels per call

Both arms keep the same explicit inner graph break. Region boxes summarize actual compiler output, ordered by AOT compile ID.

This overview does not assert CUDA Graph dependency edges. Inspect the per-region operator DAGs and recorded replay timeline in the interactive view.