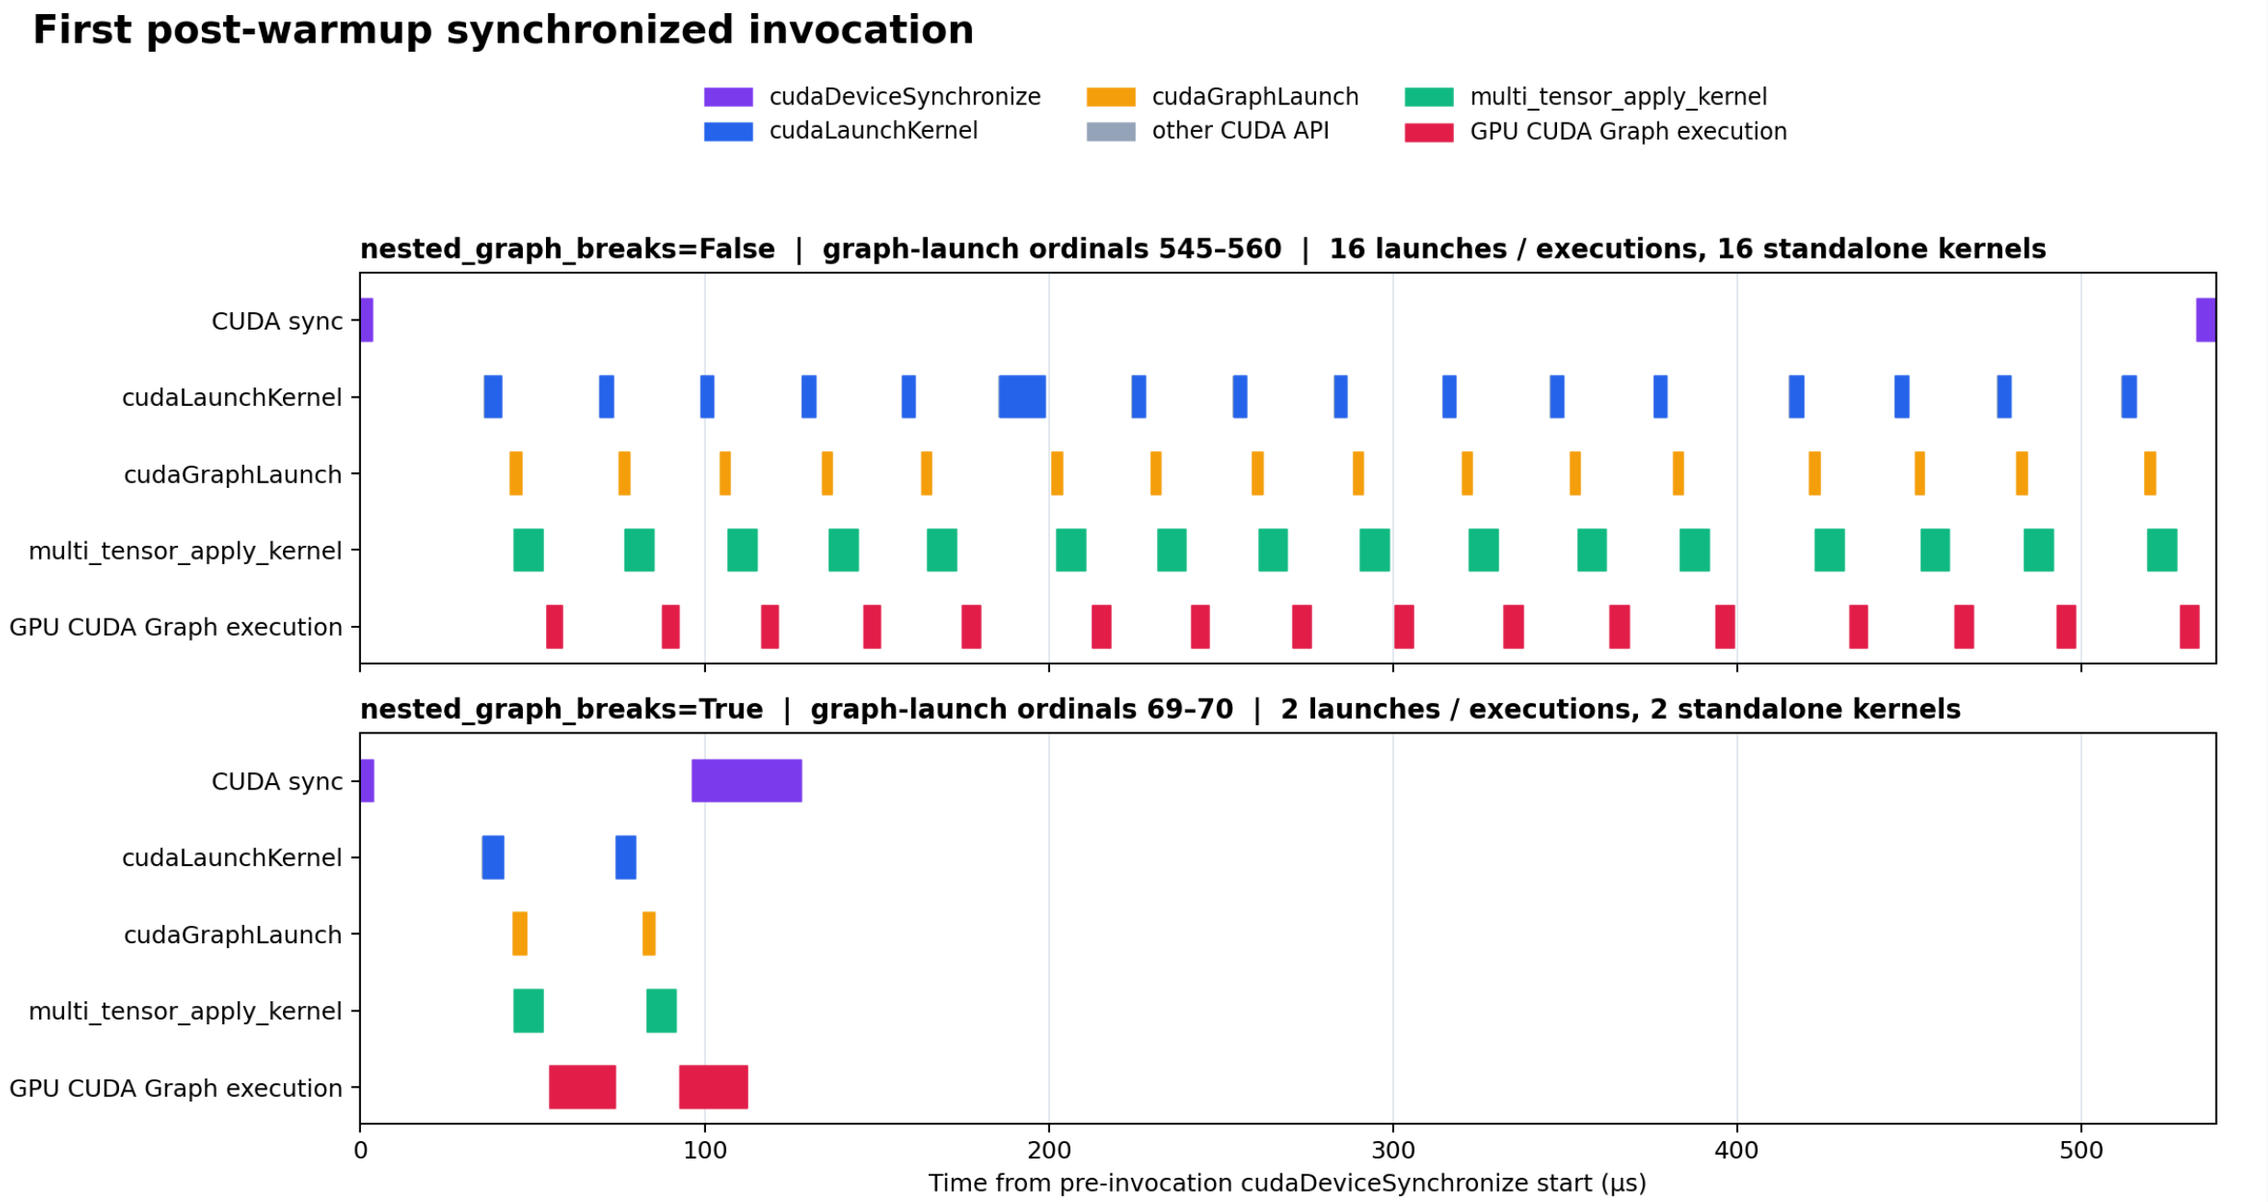
RETAINED NVIDIA PROFILE

Nsys: first post-warmup call

B200 · PyTorch 2.14.1 · graph-level trace

Top: without nested graph-break support
Bottom: with nested graph-break support



Rendered from retained Nsys SQLite in graph-level trace mode. Shared time scale and exact spans. No clipping, smoothing, interpolation, or CPU/GPU causal fraction.

One synchronized invocation per arm.

Retained timeline fitted in full. Shared time axis in *µ*s.

Graph-level trace, separate from acceptance timing.

Original capture, report and code linked with the post.

RETAINED NVIDIA PROFILE

NCU: one kernel record per arm

B200 · retained Nsight Compute metrics

False column: without nested graph-break support
True column: with nested graph-break support

Nsight Compute: one retained kernel record per arm

Kernel: triton_poi_fused_add_cos_mul_sin_0 | Device: NVIDIA B200

Raw CSV column	Unit	False	True
gpu__time_duration.avg	us	6.400000	20.992000
gpu__dram_throughput.avg.pct_of_peak_sustained_elapsed	%	17.261867	5.248152
sm__throughput.avg.pct_of_peak_sustained_elapsed	%	30.132462	64.497376
sm__warps_active.avg.pct_of_peak_sustained_active	%	75.796713	82.802779
lts__throughput.avg.pct_of_peak_sustained_elapsed	%	11.632164	3.495593
launch__grid_size		2048	2048
launch__block_size		256	256
launch__registers_per_thread	register/thread	18	18
launch__shared_mem_per_block	Kbyte/block	1.024000	1.024000
profiler__replayer_passes	pass	9.000000	9.000000

Values and units are copied verbatim from the two retained wide raw-page CSV rows. No ratio or aggregate is calculated.

Single-kernel profiler replay.

Retained report values and units remain unchanged.

Separate from unprofiled acceptance timing.

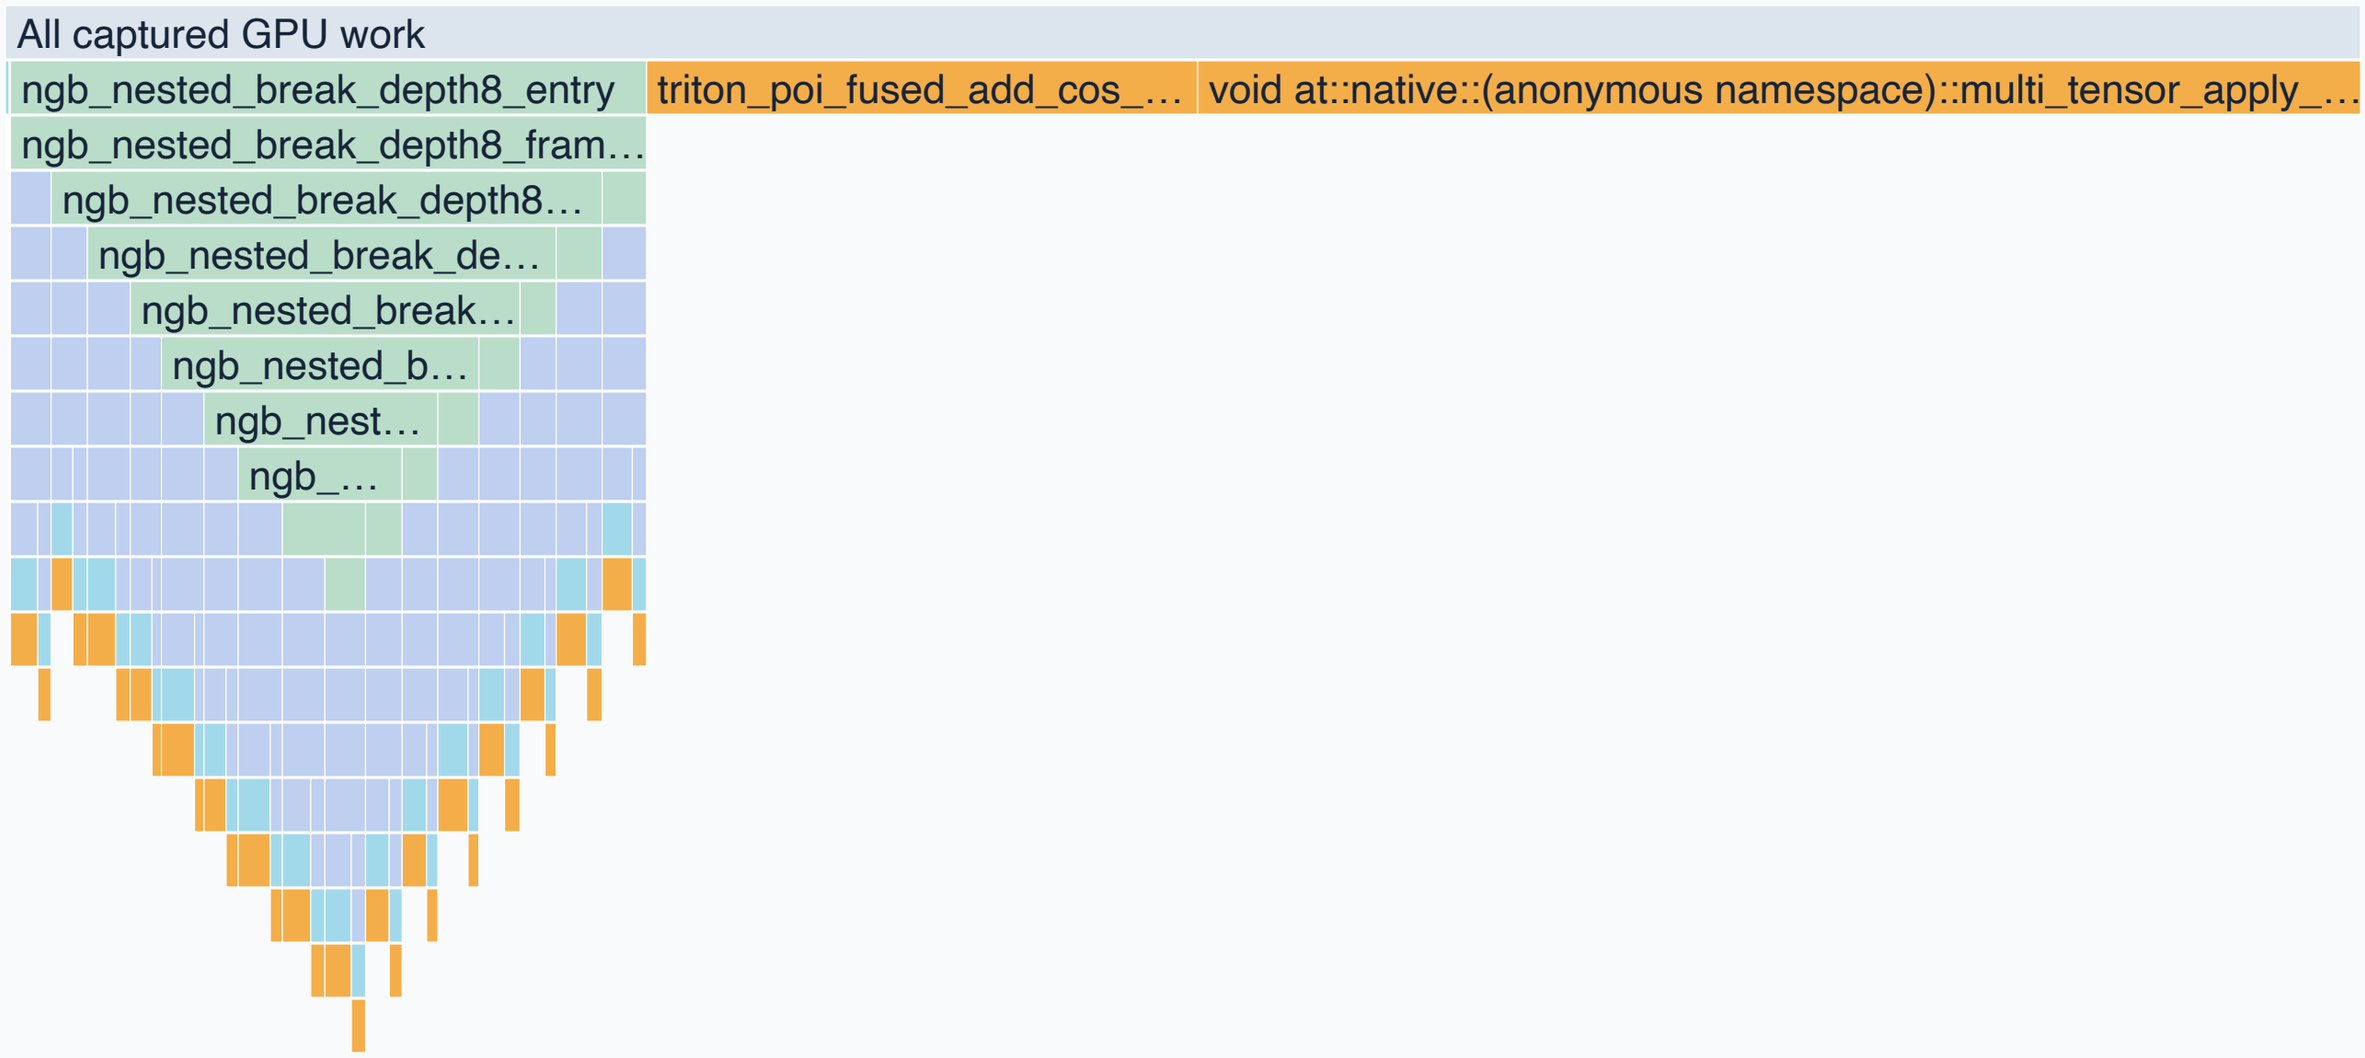
Original capture, report and code linked with the post.

ZYMTRACE / CAPTURED GPU TIME

Without nested graph-break support

B200 · PyTorch 2.14.1 · depth-8 microbenchmark

Bar width = % of this worker's captured GPU time



Key functions · inclusive GPU time

CUDAGraphNode.run	**27.01%**
CUDAGraph.replay	**8.73%**
cuLaunchKernel	**18.45%**
cuGraphLaunch	**8.73%**

Each arm normalized separately. Inclusive shares overlap.

Includes setup, warmup, timing and correctness work.

GPU time attributed to host frames is not CPU execution time.

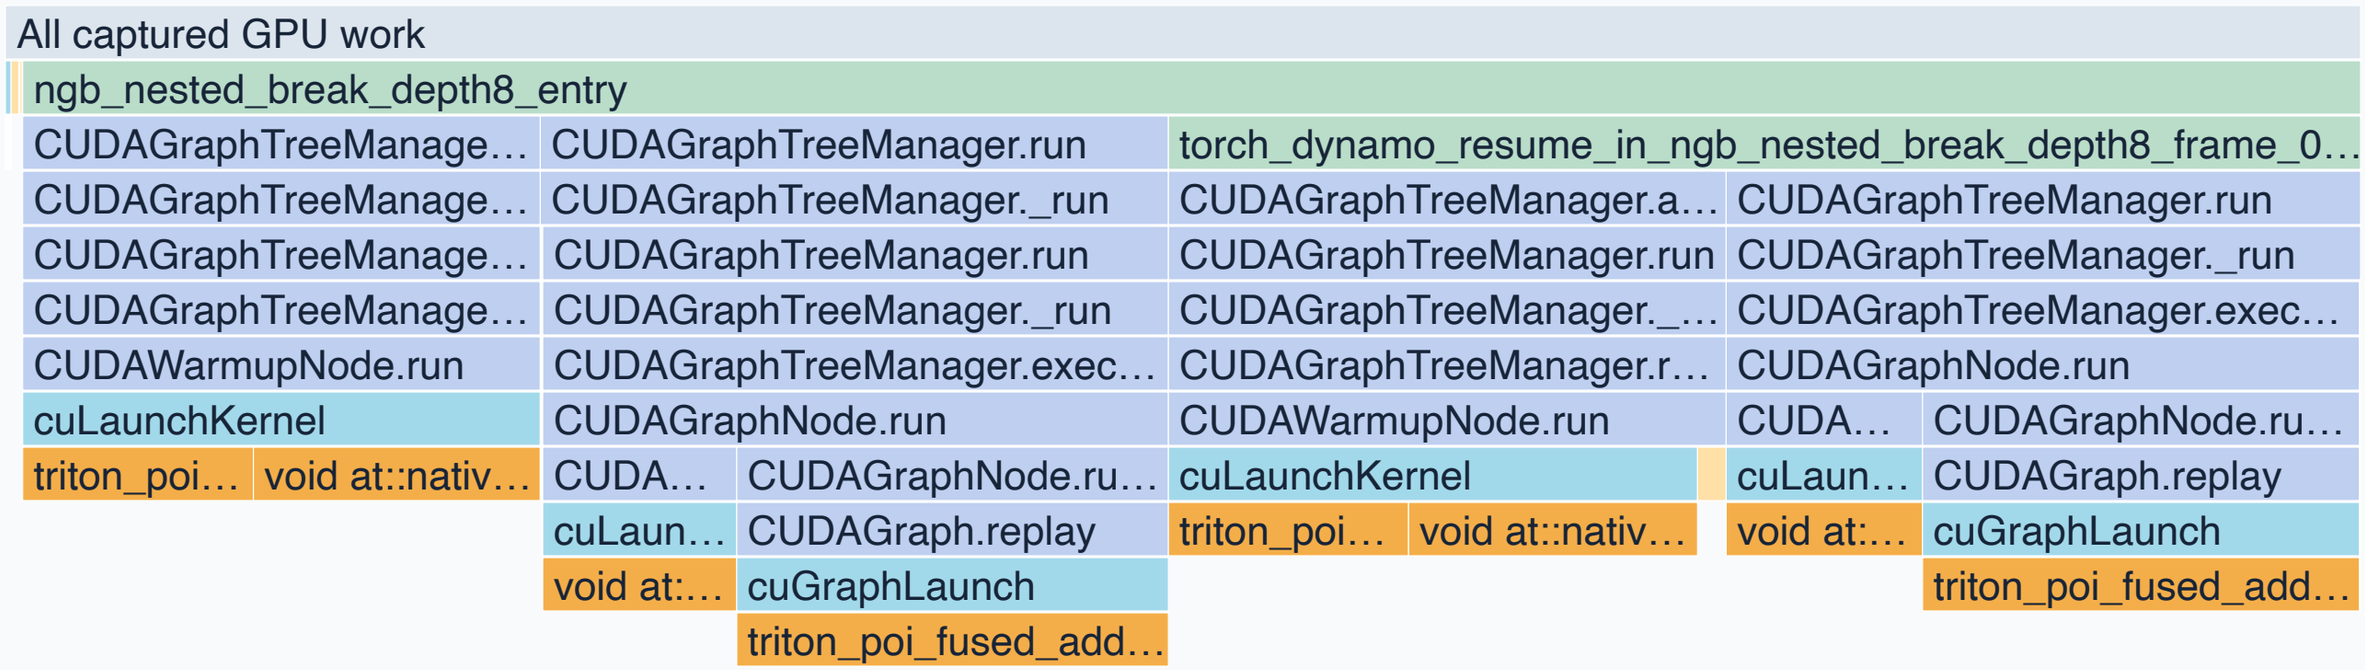
Not steady-state latency. Full function stacks linked with the post.

ZYMTRACE / CAPTURED GPU TIME

With nested graph-break support

B200 · PyTorch 2.14.1 · depth-8 microbenchmark

Bar width = % of this worker's captured GPU time



Key functions · inclusive GPU time

CUDAGraphNode.run	**53.50%**
CUDAGraph.replay	**36.98%**
cuLaunchKernel	**61.31%**
cuGraphLaunch	**36.98%**

Each arm normalized separately. Inclusive shares overlap.

Includes setup, warmup, timing and correctness work.

GPU time attributed to host frames is not CPU execution time.

Not steady-state latency. Full function stacks linked with the post.